

Name: Parsa Alamzadeh
Computing ID: palamzad
Kaggle Team Name: **Papa Goose**

Date: Nov. 1st, 2020
Student number: 301316272

Project 3: Semantic Segmentation

Link to my colab notebook:

<https://colab.research.google.com/drive/1cxwu0Fm2cvCTd4Mz-J5Td4SQWm9v8eq0?usp=sharing>

Part 1:

For this section, I first cached the data into a variable so I don't have to read from the disk every time. This significantly reduced the time of training, 2-3 mins preloading as opposed to 2-3 minutes every time for training.

For configuration of the pretrained model I only modified the MAX_ITER and the BASE_LR. Here is the configuration I used that outperformed the other configurations:


BASE_LR: 0.00025

MAX_ITER: 1,000


Different learning rates converged to their final accuracy in different rates. The one with 0.001 started off with around 65% accuracy and barely improved (converged around 72-73%). 0.0005 also had a similar issue where it started off with a high accuracy but failed to improve as much, therefore it was still a bit high and when I chose 0.00025 I saw that accuracy steadily improved. A higher BASE_LR had a lower accuracy while a lower learning struggled to converge. This also can be seen in the Total Loss graph as well. The reason for that is that too large of a learning rate cannot converge to local minima while too small of a learning rate can't find the global minima and gets stuck at the local minima.

Another modification that I made was using the [faster_rcnn_X_101_32x8d_FPN_3x](#).yaml instead of the suggested one, since it had a better convergence compared to the provided one.

Here are the graphs from different configurations:
Accuracy of the models:



Total loss:



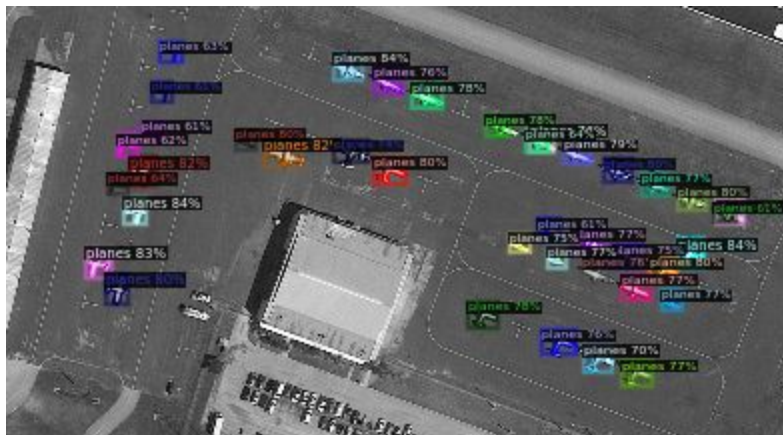
The straight lines are the linear regression automatically generated which I couldn't remove or modify. The learning rates I tried were: 0.001, 0.0005, 0.00025 and 0.0001. The reason I used a higher MAX_ITER was to allow the lower learning rate to converge.

Here are some visualizations done on the test dataset.
Visualizations from the optimal model:





The results from a higher learning rate:





While the masks are not bad, it is missing a lot of planes.

Here is the results from a lower learning rate:







This model fails completely to make any predictions with a high confidence.

As for evaluation, I set aside 10 (~5% of the data) images from the training dataset for validation. Here is the results from COCOevaluator on the validation data:

BBox results:


```
Evaluate annotation type *bbox*
COCOeval_opt.evaluate() finished in 0.02 seconds.
Accumulating evaluation results...
COCOeval_opt.accumulate() finished in 0.00 seconds.
```

Average Precision	(AP)	@[IoU=0.50:0.95	area= all	maxDets=100] = 0.312
Average Precision	(AP)	@[IoU=0.50	area= all	maxDets=100] = 0.550
Average Precision	(AP)	@[IoU=0.75	area= all	maxDets=100] = 0.329
Average Precision	(AP)	@[IoU=0.50:0.95	area= small	maxDets=100] = 0.359
Average Precision	(AP)	@[IoU=0.50:0.95	area=medium	maxDets=100] = 0.227
Average Precision	(AP)	@[IoU=0.50:0.95	area= large	maxDets=100] = 0.735
Average Recall	(AR)	@[IoU=0.50:0.95	area= all	maxDets= 1] = 0.014
Average Recall	(AR)	@[IoU=0.50:0.95	area= all	maxDets= 10] = 0.109
Average Recall	(AR)	@[IoU=0.50:0.95	area= all	maxDets=100] = 0.338
Average Recall	(AR)	@[IoU=0.50:0.95	area= small	maxDets=100] = 0.367
Average Recall	(AR)	@[IoU=0.50:0.95	area=medium	maxDets=100] = 0.245
Average Recall	(AR)	@[IoU=0.50:0.95	area= large	maxDets=100] = 0.829

```
[11/01 08:18:20 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
```

AP	AP50	AP75	APs	APm	APl
31.216	54.958	32.930	35.947	22.717	73.536

Segmentation results:

Average Precision	(AP)	@[IoU=0.50:0.95	area= all	maxDets=100] = 0.113
Average Precision	(AP)	@[IoU=0.50	area= all	maxDets=100] = 0.385
Average Precision	(AP)	@[IoU=0.75	area= all	maxDets=100] = 0.019
Average Precision	(AP)	@[IoU=0.50:0.95	area= small	maxDets=100] = 0.099
Average Precision	(AP)	@[IoU=0.50:0.95	area=medium	maxDets=100] = 0.088
Average Precision	(AP)	@[IoU=0.50:0.95	area= large	maxDets=100] = 0.447
Average Recall	(AR)	@[IoU=0.50:0.95	area= all	maxDets= 1] = 0.007
Average Recall	(AR)	@[IoU=0.50:0.95	area= all	maxDets= 10] = 0.054
Average Recall	(AR)	@[IoU=0.50:0.95	area= all	maxDets=100] = 0.142
Average Recall	(AR)	@[IoU=0.50:0.95	area= small	maxDets=100] = 0.128
Average Recall	(AR)	@[IoU=0.50:0.95	area=medium	maxDets=100] = 0.104
Average Recall	(AR)	@[IoU=0.50:0.95	area= large	maxDets=100] = 0.497

```
[11/01 08:18:20 d2.evaluation.coco_evaluation]: Evaluation results for segm:
```

AP	AP50	AP75	APs	APm	APl
11.300	38.490	1.922	9.876	8.814	44.710

Part 2:

Similar to part 1, I cached the data needed for training the model. As for the hyperparameters, I had the epochs set to 50, with batch size of 32 with initial learning rate of 0.1 and weight decay of 0.00001. Every 15 epochs (after the 30th epoch), I decreased the learning rate by a factor of 10.


Here is the architecture of the network I used for my predictions:

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 8, 128, 128]	224
BatchNorm2d-2	[-1, 8, 128, 128]	16
ReLU-3	[-1, 8, 128, 128]	0
conv-4	[-1, 8, 128, 128]	0
Conv2d-5	[-1, 16, 128, 128]	1,168
BatchNorm2d-6	[-1, 16, 128, 128]	32
ReLU-7	[-1, 16, 128, 128]	0
conv-8	[-1, 16, 128, 128]	0
MaxPool2d-9	[-1, 16, 64, 64]	0
down-10	[-1, 16, 64, 64]	0
Conv2d-11	[-1, 32, 64, 64]	4,640
BatchNorm2d-12	[-1, 32, 64, 64]	64
ReLU-13	[-1, 32, 64, 64]	0
conv-14	[-1, 32, 64, 64]	0
MaxPool2d-15	[-1, 32, 32, 32]	0
down-16	[-1, 32, 32, 32]	0
Conv2d-17	[-1, 64, 32, 32]	18,496
BatchNorm2d-18	[-1, 64, 32, 32]	128
ReLU-19	[-1, 64, 32, 32]	0
conv-20	[-1, 64, 32, 32]	0
MaxPool2d-21	[-1, 64, 16, 16]	0
down-22	[-1, 64, 16, 16]	0
Conv2d-23	[-1, 128, 16, 16]	73,856
BatchNorm2d-24	[-1, 128, 16, 16]	256
ReLU-25	[-1, 128, 16, 16]	0
conv-26	[-1, 128, 16, 16]	0
MaxPool2d-27	[-1, 128, 8, 8]	0
down-28	[-1, 128, 8, 8]	0
Conv2d-29	[-1, 256, 8, 8]	295,168
BatchNorm2d-30	[-1, 256, 8, 8]	512
ReLU-31	[-1, 256, 8, 8]	0
conv-32	[-1, 256, 8, 8]	0
MaxPool2d-33	[-1, 256, 4, 4]	0
down-34	[-1, 256, 4, 4]	0
Conv2d-35	[-1, 512, 4, 4]	1,180,160
BatchNorm2d-36	[-1, 512, 4, 4]	1,024
ReLU-37	[-1, 512, 4, 4]	0
conv-38	[-1, 512, 4, 4]	0
MaxPool2d-39	[-1, 512, 2, 2]	0
down-40	[-1, 512, 2, 2]	0
ConvTranspose2d-41	[-1, 512, 4, 4]	1,049,088
Conv2d-42	[-1, 256, 4, 4]	1,179,904
BatchNorm2d-43	[-1, 256, 4, 4]	512
ReLU-44	[-1, 256, 4, 4]	0
conv-45	[-1, 256, 4, 4]	0
up-46	[-1, 256, 4, 4]	0
BatchNorm2d-47	[-1, 256, 4, 4]	512
ConvTranspose2d-48	[-1, 256, 8, 8]	262,400
Conv2d-49	[-1, 128, 8, 8]	295,040
BatchNorm2d-50	[-1, 128, 8, 8]	256
ReLU-51	[-1, 128, 8, 8]	0
conv-52	[-1, 128, 8, 8]	0
up-53	[-1, 128, 8, 8]	0
BatchNorm2d-54	[-1, 128, 8, 8]	256
ConvTranspose2d-55	[-1, 128, 16, 16]	65,664
Conv2d-56	[-1, 64, 16, 16]	73,792
BatchNorm2d-57	[-1, 64, 16, 16]	128
ReLU-58	[-1, 64, 16, 16]	0
conv-59	[-1, 64, 16, 16]	0
up-60	[-1, 64, 16, 16]	0
BatchNorm2d-61	[-1, 64, 16, 16]	128
ConvTranspose2d-62	[-1, 64, 32, 32]	16,448
Conv2d-63	[-1, 32, 32, 32]	18,464
BatchNorm2d-64	[-1, 32, 32, 32]	64
ReLU-65	[-1, 32, 32, 32]	0
conv-66	[-1, 32, 32, 32]	0
up-67	[-1, 32, 32, 32]	0
BatchNorm2d-68	[-1, 32, 32, 32]	64
ConvTranspose2d-69	[-1, 32, 64, 64]	4,128
Conv2d-70	[-1, 16, 64, 64]	4,624
BatchNorm2d-71	[-1, 16, 64, 64]	32
ReLU-72	[-1, 16, 64, 64]	0
conv-73	[-1, 16, 64, 64]	0
up-74	[-1, 16, 64, 64]	0
BatchNorm2d-75	[-1, 16, 64, 64]	32
ConvTranspose2d-76	[-1, 16, 128, 128]	1,040
Conv2d-77	[-1, 8, 128, 128]	1,160
BatchNorm2d-78	[-1, 8, 128, 128]	16
ReLU-79	[-1, 8, 128, 128]	0
conv-80	[-1, 8, 128, 128]	0
up-81	[-1, 8, 128, 128]	0
BatchNorm2d-82	[-1, 8, 128, 128]	16
Conv2d-83	[-1, 3, 128, 128]	219
BatchNorm2d-84	[-1, 3, 128, 128]	6
ReLU-85	[-1, 3, 128, 128]	0
conv-86	[-1, 3, 128, 128]	0
Conv2d-87	[-1, 1, 128, 128]	28
conv-88	[-1, 1, 128, 128]	0

The model's architecture is very similar to the existing one, but I also borrowed the skip layer component from resnet, which improved the accuracy of the model. Output of downsampling layers are copied and added to the upsampling ones. I have 6 downsampling layers and 6 up sampling. As for the optimizer and loss function, I used the default ones and achieved loss of 0.07.



For IoU, I did an element-wise multiplication and counted the non-zero elements (for the intersection operation) . For the union I did element-wise addition and counted the non-zero elements. This is after setting all pixel values to 0 and 1 for both the predicted and ground truth. The average IoU for all instances was 0.90.



Below are samples of my models predictions on the training dataset, where the top image is the actual image, the second one is model's prediction and the third one is the ground truth mask:



Part 3:

I submitted my results as **Papa Goose** on kaggle. My current score is 0.62312 on the public leaderboard. Here are some of the predicted masks from my model on the training dataset, the first one is the actual image, second one is the ground truth mask and the third one is the predicted mask:





Part 4:

Visualizations:



This model performs way worse than the one used in part 1. Here are the AP scores for this model on the validation data:

BBox AP:

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.236
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.476
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.211
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.204
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.272
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.500
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.014
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.110
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.292
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.215
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.336
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.649

```

[11/02 02:27:03 d2.evaluation.coco_evaluation]: Evaluation results for bbox:

AP	AP50	AP75	APs	APm	APl
23.555	47.633	21.108	20.375	27.195	50.033

Segmentation AP:

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.042
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.182
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.007
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.020
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.057
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.141
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.007
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.043
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.074
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.042
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.095
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.198

```

[11/02 02:27:04 d2.evaluation.coco_evaluation]: Evaluation results for segm:

AP	AP50	AP75	APs	APm	APl
4.175	18.161	0.721	2.049	5.745	14.067